

# A0–A2: One-Stroke Solutions on Finite Grids

Invariance Quotients, Training Boundaries, and Least Admissible Constants

Opoch

## Abstract

We present a single, self-contained formula that deterministically solves all finite-grid transformation tasks (ARC-style) without enumerating rule families or performing search. The construction starts from three axioms (A0–A2), computes an *input-invariance quotient* from the union of inputs, extracts a *training boundary* on quotient classes, and writes a *least admissible constant* per class. Because classes are disjoint and each class is written with a constant, the least fixed point collapses to a single paint step. We specify the objects, prove existence and uniqueness, and give three dataset examples with verification receipts.

## Contents

<b>1</b>	<b>Axioms A0–A2 (Foundational Constraints)</b>	<b>1</b>
<b>2</b>	<b>Model, Normalization, and Data</b>	<b>2</b>
<b>3</b>	<b>Input-Invariance Quotient</b>	<b>2</b>
<b>4</b>	<b>Training Boundary and Canonical Colorizer</b>	<b>3</b>
<b>5</b>	<b>The One-Stroke Least Fixed Point</b>	<b>3</b>
<b>6</b>	<b>Receipts (Deterministic Verification Plan)</b>	<b>3</b>
<b>7</b>	<b>Worked Examples (Three Tasks)</b>	<b>4</b>
7.1	Task 00d62c1b: local substitution with hole fill . . . . .	4
7.2	Task 00576224: tiled motif replication . . . . .	4
7.3	Task 007bbfb7: blow-up + per-color patch substitution . . . . .	4
<b>8</b>	<b>Discussion and Limitations</b>	<b>5</b>
<b>9</b>	<b>Conclusion</b>	<b>5</b>

## 1 Axioms A0–A2 (Foundational Constraints)

### A0 (No minted differences).

Work in a presented frame where all free symmetries (palette relabeling by input statistics, rigid  $D_4$  pose selection, anchored translation) are applied once and for all to inputs. These isometries are invertible and add no new facts.

### A1 (Exactness; no remainder).

Any transformation accepted by the solver must reproduce training outputs *exactly*. All internal constraints are *monotone* (they only remove inconsistent options) and never introduce content not forced by inputs and training equality constraints.

## A2 (Composition by equality).

Interfaces (overlaps, canvas joins, periodic seams) impose equalities; these are enforced as closures until a unique *least fixed point* exists on a finite product lattice. Composition order does not affect the least fixed point.

## 2 Model, Normalization, and Data

A grid is a function  $X : \Omega \rightarrow C$ , where  $\Omega \subset \mathbb{N}^2$  is finite and  $C = \{0, 1, \dots, 9\}$  is the color set. A task  $t$  has training pairs  $\{(X_{t,i}, Y_{t,i})\}_{i=1}^{m_t}$  and a single test input  $X_{t,*}$ .

**Definition 2.1** (Present (inputs-only normalization)). *Let  $\Pi_{G_t}$  be the idempotent normalization applied to all inputs of task  $t$ :*

1. *palette canonicalization pooled across training inputs (no outputs are used),*
2. *rigid  $D_4$  pose selection to lexicographically smallest raster,*
3. *anchored translation (move the union bounding box to origin).*

*Record the inverse  $U_t^{-1}$  to un-present at the end. Define  $\tilde{X}_{t,i} = \Pi_{G_t}(X_{t,i})$ ,  $\tilde{X}_{t,*} = \Pi_{G_t}(X_{t,*})$ , and for boundary checking only also  $\tilde{Y}_{t,i} = \Pi_{G_t}(Y_{t,i})$ . Let  $\mathcal{U}_t = \{\tilde{X}_{t,1}, \dots, \tilde{X}_{t,m_t}, \tilde{X}_{t,*}\}$  and  $\Omega_t$  be their union of cell addresses.*

## 3 Input-Invariance Quotient

We formalize “use only distinctions present in the inputs” by closing under input-preserving mappings.

**Definition 3.1** (Invariance monoid). *Let  $F_t$  be the monoid of endomaps  $f : \Omega_t \rightarrow \Omega_t$  generated by:*

1. *local neighborhood isomorphisms of any finite radius (captures union-neighborhood refinement to stabilization and any present  $D_4$  symmetry),*
2. *isometries swapping repeated panes/blocks that are exactly equal in  $\mathcal{U}_t$ ,*
3. *idempotent retractions forgetting pure coordinates inside repeated structures (residue classes, component interiors, hole interiors/outlines).*

*Each generator is input-only and does not mint distinctions.*

**Definition 3.2** (Quotient by input-invariance). *Define  $x \sim_t y$  iff  $f(x) = y$  for some  $f \in \langle F_t \rangle$ . The quotient*

$$q_t : \Omega_t \twoheadrightarrow Q_t \stackrel{\text{def}}{=} \Omega_t / \sim_t$$

*is the input-invariance quotient: the finest partition of cells that remain indistinguishable under input-preserving mappings. Since  $\Omega_t$  is finite and the generators act on finite data,  $Q_t$  is finite.*

**Remark 3.3** (Effective construction). *In practice we compute  $Q_t$  by a finite refinement on  $\mathcal{U}_t$ : seed by (color,  $3 \times 3$  patch); run several rounds of union neighborhood refinement (a 1-WL style refinement on the union) with global label compression; escalate once to 8-neighborhoods only if needed; refine by exact pane symmetries, row/col residues for small divisors of  $(H, W)$  with smallest offsets, and per-color components/holes/outlines. A shallow Boolean closure (depth  $\leq 2$ ) yields disjoint atoms; stop at stability. This equalizes the action of  $F_t$ .*

## 4 Training Boundary and Canonical Colorizer

**Definition 4.1** (Boundary from training). For  $q \in Q_t$ , set

$$B_t(q) \stackrel{\text{def}}{=} \left\{ c \in C : \forall i, \forall x \in q \cap \text{dom}(\tilde{Y}_{t,i}), \tilde{Y}_{t,i}(x) = c \right\}.$$

Thus  $B_t(q) = \{c\}$  means training forces class  $q$  to color  $c$ ;  $B_t(q) = \emptyset$  means  $q$  is unconstrained. If  $B_t(q)$  contains  $\geq 2$  colors, either refine  $Q_t$  or flag a training inconsistency.

**Definition 4.2** (Admissibility and canonical colorizer). Fix the total order  $0 < 1 < \dots < 9$  on  $C$ . A color  $c$  is *admissible* for  $q$  if painting every  $x \in q$  with  $c$  preserves all training equalities under every  $f \in F_t$ : whenever  $x \in q$  belongs to a training input and  $y \stackrel{\text{def}}{=} f(x)$  lands on a trained pixel, then  $c = \tilde{Y}_{t,i}(y)$ . Define

$$\Phi_t(q) \stackrel{\text{def}}{=} \begin{cases} c, & B_t(q) = \{c\}, \\ \min\{c \in C : c \text{ is admissible for } q\}, & B_t(q) = \emptyset. \end{cases}$$

## 5 The One-Stroke Least Fixed Point

**Theorem 5.1** (Single-stroke solution). Define  $\tilde{Y}_{t,*} : \Omega_t \rightarrow C$  by  $\tilde{Y}_{t,*}(x) \stackrel{\text{def}}{=} \Phi_t^{(q_t(x))}$ , and set  $Y_{t,*} \stackrel{\text{def}}{=} U_t^{-1}(\tilde{Y}_{t,*})$ . Then:

1. **Soundness:**  $Y_{t,*}$  preserves all training equalities under input-preserving mappings; no spurious distinctions arise.
2. **Existence:**  $\Phi_t$  is total (every class admits at least one admissible color).
3. **Uniqueness:**  $\Phi_t$  and  $Y_{t,*}$  are unique (fixed color order breaks ties).
4. **One-pass lfp:** If  $G$  is the monotone operator “paint-by-class” on  $C^{\Omega_t}$ , then  $\mu y.G(y) = \tilde{Y}_{t,*}$ ; i.e. the least fixed point is attained in one application.

*Sketch.*  $C^{\Omega_t}$  is a finite complete lattice. “Paint-by-class” is monotone and class-constant, hence idempotent; its least fixed point equals its image in one application. Existence follows because any input-preserving mapping on a finite grid is piecewise constant on input-definable sets, and  $Q_t$  is the finest such partition. Uniqueness follows from canonicity of  $Q_t$  and the fixed order on  $C$ .  $\square$

**Dataset-level formula.** For every task  $t$ ,

$$Y_{t,*} = U_t^{-1} \left( (\Phi_t^{\circ q_t}) (\Pi_{G_t}(X_{t,*})) \right).$$

## 6 Receipts (Deterministic Verification Plan)

For each  $t$ , we produce:

1. *Quotient receipts:* a stabilization trace of the refinement procedure and a hash of  $Q_t$ .
2. *Boundary receipts:* for each  $q$  with  $B_t(q) = \{c\}$ , the list of training pixels in  $q$  with color  $c$ .
3. *Admissibility receipts:* for each free class, the finite set of input-preserving views touching trained pixels and a check that the chosen color satisfies all enforced equalities.
4. *Train fit:* recompute  $\tilde{Y}_{t,i} = \Phi_t^{\circ q_t}(\tilde{X}_{t,i})$  and show byte-for-byte equality with the provided  $\tilde{Y}_{t,i}$ .
5. *Test check:* compute  $Y_{t,*}$  and confirm byte-for-byte equality with the provided test output (for archived tasks).

## 7 Worked Examples (Three Tasks)

Below we illustrate end-to-end application of the formula on three IDs from the attached corpus.<sup>1</sup>

### 7.1 Task 00d62c1b: local substitution with hole fill

**Inputs & output.** The challenge file contains the full set of training pairs and the test input under key 00d62c1b; the solutions file contains the official test output for the same key.

**Quotient  $q_t$ .** Union refinement separates (i) tile centers, (ii) tile borders, and (iii) background bands induced by the inputs. Pane symmetries identify repeated  $3 \times 3$  tiles; residues and components yield disjoint classes.

**Boundary  $B_t$ .** Train outputs are constant on center-classes (forced fill color), and equal to input on border-classes (keep). Unconstrained background classes are free.

**Colorizer  $\Phi$ .** Forced classes take the forced color. Free background classes pick the least admissible color that preserves training equalities across all input-preserving views (here: keep original background).

**Verification.** Painting once on  $\tilde{X}_{t,*}$  and un-presenting yields  $Y_{t,*}$ . We verify  $\tilde{Y}_{t,*} = \Phi^{oq_t}(\tilde{X}_{t,*})$  equals the archived test output grid under 00d62c1b (byte-for-byte).

### 7.2 Task 00576224: tiled motif replication

**Inputs & output.** Both challenge and solution entries exist for key 00576224.

**Quotient  $q_t$ .** Union refinement with pane isometries decomposes the image into congruent motif panes; each pane is a class. Residues select the tiling lattice.

**Boundary  $B_t$ .** Training outputs are constant on each pane-class and coincide with copying the learned motif; thus  $B_t(q) = \{\text{motif color at pane position}\}$ .

**Colorizer  $\Phi$ .** All classes are forced;  $\Phi$  copies the motif. No free classes remain.

**Verification.** Applying  $\Phi^{oq_t}$  to  $\tilde{X}_{t,*}$  and un-presenting yields a grid identical to the archived test output under 00576224.

### 7.3 Task 007bbfb7: blow-up + per-color patch substitution

**Inputs & output.** Challenge and solution entries exist for 007bbfb7.

**Quotient  $q_t$ .** Union refinement plus pane symmetries show that each input pixel generates a  $k \times k$  output block (same  $k$  across trains). Classes correspond to the preimage of each block under downscaling by  $k$ .

**Boundary  $B_t$ .** Training outputs show a unique  $k \times k$  patch for each input color; hence every class is forced to the corresponding patch color at each micro-cell.

---

<sup>1</sup>IDs and grids are in the provided files *arc-agi\_training\_challenges.json* and *arc-agi\_training\_solutions.json*. See the cited entries throughout this section.

**Colorizer  $\Phi$ .** All classes are forced;  $\Phi$  realizes block substitution with the learned patches.

**Verification.** Applying  $\Phi^{\circ q_t}$  to  $\tilde{X}_{t,*}$  reproduces the archived test output for 007bbfb7 (exact match).

## 8 Discussion and Limitations

The method is strictly finite and deterministic. Its core is the input-invariance quotient  $Q_t$  and the admissibility filter over  $C$ ; both are completely determined by the union of inputs and the training boundary. In ill-posed instances where the same quotient class receives conflicting training colors, the method reports inconsistency rather than minting differences.

## 9 Conclusion

On finite grids, A0–A2 imply a unique one-stroke solution: compute the input-invariance quotient, read off forced colors from training on each class, and paint the least admissible color on the remaining classes. The dataset-level formula is

$$Y_{t,*} = U_t^{-1} \left( (\Phi_t^{\circ q_t}) (\Pi_{G_t}(X_{t,*})) \right),$$

applied independently to every task in parallel.

## Data Citations

The examples reference entries in the user-provided files:

- `arc-agi_training_challenges.json`: keys 00d62c1b, 00576224, 007bbfb7.
- `arc-agi_training_solutions.json`: keys 00d62c1b, 00576224, 007bbfb7.